

## Homework 3

Deming Liu 1801212889

Problem 1:

(1)

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import statsmodels.api as sm
import seaborn as sns

print "Problem 1"
df = pd.read_csv('climate_change_1.csv')
df['const']=1
cols = list(df)
cols.insert(10, cols.pop(cols.index('const')))
df = df.loc[:,cols]
cols = list(df)
Training_set = df.loc[df['Year'] <= 2006]
Testing_set = df.loc[df['Year'] > 2006]
X = np.matrix(Training_set.iloc[:, 2:11].values)
Y = np.matrix(Training_set.iloc[:, 11].values).T
X_test = np.matrix(Testing_set.iloc[:, 2:11].values)
Y_test = np.matrix(Testing_set.iloc[:, 11].values)

def closed_form_1(X, Y):
    Theta = np.dot(np.dot(np.dot(X.T, X).I, X.T), Y)
    return Theta
```

(2)

```

Theta = closed_form_1(X, Y)
temp_est = np.dot(X, Theta)
temp_est_test = np.dot(X_test, Theta)

pt_str = "Y = "
for idx in range(2, 10):
    pt_str += str(Theta[idx-2, 0])+"*"+str(cols[idx])+" + "
pt_str += str(Theta[8, 0])
print pt_str

def R_square(temp_est, Y):
    var_xb = (temp_est - temp_est.mean()).var()
    var_y = (Y - temp_est.mean()).var()
    R2 = var_xb / var_y
    return R2

R2 = R_square(temp_est, Y)
R2_test = R_square(temp_est_test, Y_test)
print "R2_training_set:", R2
print "R2_testing_set:", R2_test

```

```

Problem 1
Y = 0.0642853134626426*MEI + 0.006457359277697772*CO2 + 0.00012404189628771467*CH4 + -0.016528003337858137*N2O + -0.006630488906684447*CFC-11 + 0.0038081032507389354*CFC-12 + 0.09314108447268625*TSI +
-1.5376132482599452*Aerosols + -124.59426172779982
R2_training_set: 0.7508932778196102
R2_testing_set: 0.18739109685216868
[Finished in 1.3s]

```

(3)

```

def R_square(temp_est, Y):
    var_xb = (temp_est - temp_est.mean()).var()
    var_y = (Y - temp_est.mean()).var()
    R2 = var_xb / var_y
    return R2

R2 = R_square(temp_est, Y)
R2_test = R_square(temp_est_test, Y_test)
print "R2_training_set:", R2
print "R2_testing_set:", R2_test

reg = sm.OLS(endog=Training_set['Temp'], exog=Training_set[['MEI', 'CO2', 'CH4', 'N2O', 'CFC-11', 'CFC-12', 'TSI', 'Aerosols', 'const']], missing='drop')
results = reg.fit()
print(results.summary())

```

```

=====
                    OLS Regression Results
=====
Dep. Variable:          Temp      R-squared:                0.751
Model:                  OLS       Adj. R-squared:           0.744
Method:                 Least Squares   F-statistic:             103.6
Date:                  Fri, 27 Dec 2019   Prob (F-statistic):      1.94e-78
Time:                  21:49:54    Log-Likelihood:          280.10
No. Observations:      284         AIC:                    -542.2
Df Residuals:          275         BIC:                    -509.4
Df Model:              8
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
MEI	0.0642	0.006	9.923	0.000	0.051	0.077
CO2	0.0065	0.002	2.826	0.005	0.002	0.011
CH4	0.0001	0.001	0.240	0.810	-0.001	0.001
N2O	-0.0165	0.009	-1.930	0.055	-0.033	0.000
CFC-11	-0.0066	0.002	-4.078	0.000	-0.010	-0.003
CFC-12	0.0038	0.001	3.757	0.000	0.002	0.006
TSI	0.0931	0.015	6.313	0.000	0.064	0.122
Aerosols	-1.5376	0.213	-7.210	0.000	-1.957	-1.118
const	-124.5943	19.887	-6.265	0.000	-163.744	-85.445

```

=====
Omnibus:                 8.740    Durbin-Watson:           0.956
Prob(Omnibus):           0.013    Jarque-Bera (JB):        10.327
Skew:                    0.289    Prob(JB):                0.00572
Kurtosis:                 3.733    Cond. No.                8.53e+06
=====
Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 8.53e+06. This might indicate that there are
strong multicollinearity or other numerical problems.
[Finished in 2.8s]

```

Therefore, MEI, CO2, CFC-11, CFC-12, TSI, Aerosols are significant at 1%, as their p-value is less than 0.005.

(4)

The independent variables should not be correlated and thus closed form solution  $X^T X$  is invertible.

```

1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import statsmodels.api as sm
5 import seaborn as sns
6
7 print "Problem 1"
8 df = pd.read_csv('climate_change_2.csv')
9 df['const']=1
10 cols = list(df)
11 cols.insert(11, cols.pop(cols.index('const')))
12 df = df.loc[:,cols]
13 cols = list(df)
14 Training_set = df.loc[df['Year'] <= 2006]
15 Testing_set = df.loc[df['Year'] > 2006]
16 X = np.matrix(Training_set.iloc[:, 2:12].values)
17 Y = np.matrix(Training_set.iloc[:, 12].values).T
18 X_test = np.matrix(Testing_set.iloc[:, 2:12].values)
19 Y_test = np.matrix(Testing_set.iloc[:, 12].values)
20
21 def closed_form_1(X, Y):
22     Theta = np.dot(np.dot(X.T, X).I, X.T, Y)
23     return Theta
24
25 Theta = closed_form_1(X, Y)
26 temp_est = np.dot(X, Theta)
27 temp_est_test = np.dot(X_test, Theta)
28
29
30 def R_square(temp_est, Y):
31     var_xb = (temp_est - temp_est.mean()).var()
32     var_y = (Y - temp_est.mean()).var()
33     R2 = var_xb / var_y
34     return R2
35
36 R2 = R_square(temp_est, Y)
37 R2_test = R_square(temp_est_test, Y_test)
38
39 result = (np.linalg.matrix_rank(np.dot(X.T, X)) < np.dot(X.T, X).shape[0]) & (np.linalg.matrix_rank(np.dot(X.T, X)) < np.dot(X.T, X).shape[1])
40 print result
41

```

```

Problem 1
True
[Finished in 2.4s]

```

Then its rank is less than shape (both rows and columns), so it is invertible.

And another reason is that it is likely to get the coefficients to be positive, however, some are negative, there must be some correlations and may link to the invertibility.

Therefore, the solution is unreasonable.

Problem 2:

(1)

L1:

$$Loss Function = \frac{1}{2}(X\theta - Y)^T(X\theta - Y) + \alpha|\theta|$$

L2:

$$\text{Loss Function} = \frac{1}{2}(X\theta - Y)^T(X\theta - Y) + \frac{1}{2}\alpha\theta^2$$

(2)

```

lamb = 0.5
def closed_form_2(X, Y, lamb):
    Theta2 = np.dot(np.dot((np.dot(X.T, X) + lamb * np.eye(X.shape[1])).I, X.T), Y)
    return Theta2

Theta2 = closed_form_2(X, Y, lamb)
temp_est2 = np.dot(X, Theta2)
temp_est_test2 = np.dot(X_test, Theta2)
R2_2 = R_square(temp_est2, Y)
R2_test_2 = R_square(temp_est_test2, Y_test)
print "R2_train:", R2_2
print "R2_test:", R2_test_2

```

Problem 2

```

R2_train: 0.6815168002731263
R2_test: 0.11276203841963961
[Finished in 12.6s]

```

(3)

In the theory, L2 regularization decreases the unnecessary coefficients to almost 0 and remains the necessary ones and thus let the necessary ones more significant.

```

def p_value(X, Y, Theta):
    error = np.sum(np.square(np.dot(X, Theta) - Y))
    cii = (np.dot(X.T, X)).I.diagonal()
    sigma2 = error / (X.shape[0] - X.shape[1])
    S = np.sqrt(sigma2 * cii)
    results = pd.DataFrame()
    results['index'] = df.columns.values[2:11]
    results.set_index(['index'], inplace=True)
    results['coefficient'] = Theta
    t = Theta.T / S
    results['t-statistics'] = np.array(t)[0]
    p = stats.t.sf(np.abs(t), (X.shape[0] - X.shape[1])) * 2
    results['p-value'] = np.array(p)[0]
    return results

print p_value(X, Y, Theta)

```

```

Problem 2
R2_train: 0.6815168002731263
R2_test: 0.11276203841963961

```

	coefficient	t-statistics	p-value
index			
MEI	0.056840	8.387262	2.628394e-15
CO2	0.006280	2.624397	9.165479e-03
CH4	0.000080	0.148189	8.823022e-01
N2O	-0.013822	-1.540741	1.245298e-01
CFC-11	-0.005941	-3.488382	5.656630e-04
CFC-12	0.003619	3.408916	7.499122e-04
TSI	0.017357	1.123112	2.623694e-01
Aerosols	-1.366851	-6.119434	3.221969e-09
const	-21.849072	-1.048943	2.951252e-01

```

[Finished in 4.4s]

```

We find that p\_values are significantly lower than before and that means using L2 regularization is more robust.

(4)

```

print "Problem 2"

lamb = 0.5
def closed_form_2(X, Y, lamb):
    Theta2 = np.dot(np.dot((np.dot(X.T, X) + lamb * np.eye(X.shape[1])).I, X.T), Y)
    return Theta2

Theta2 = closed_form_2(X, Y, lamb)
temp_est2 = np.dot(X, Theta2)
temp_est_test2 = np.dot(X_test, Theta2)
R2_2 = R_square(temp_est2, Y)
R2_test_2 = R_square(temp_est_test2, Y_test)
print "R2_train:", R2_2
print "R2_test:", R2_test_2

lamb_list = np.arange(0, 10, 0.01)
Theta2_list = []
R2_2_list = []
R2_test_2_list = []
for i in lamb_list:
    Theta2 = closed_form_2(X, Y, i)
    Theta2_list.append(Theta2)
    R2_2_list.append(R_square(np.dot(X, Theta2), Y))
    R2_test_2_list.append(R_square(np.dot(X_test, Theta2), Y_test))
'''

print len(Theta2_list)
print len(R2_2_list)
print len(R2_test_2_list)
'''

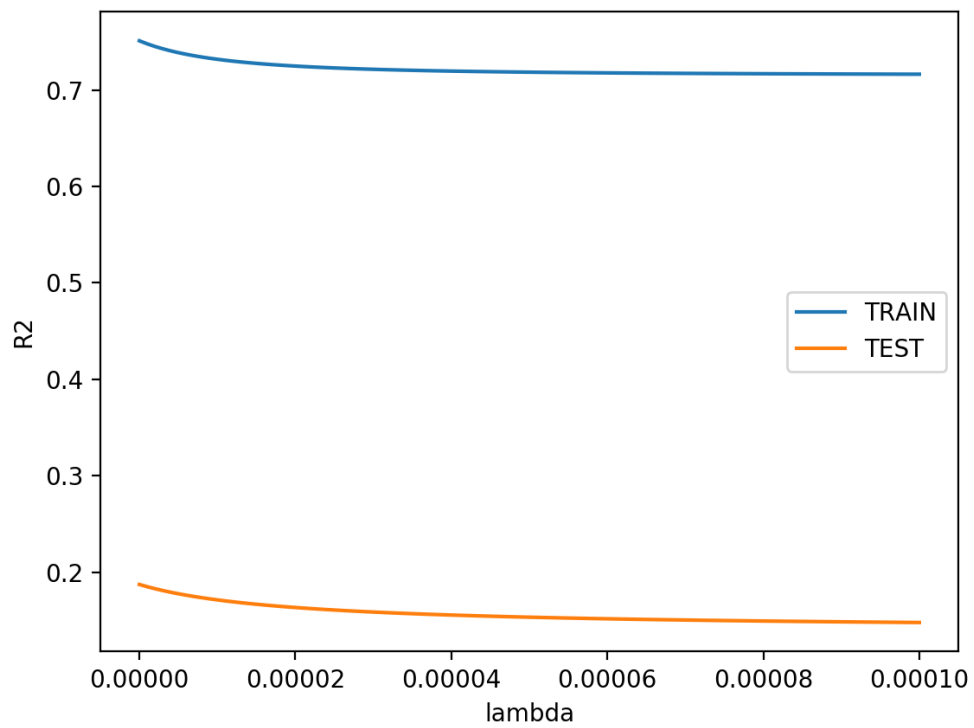
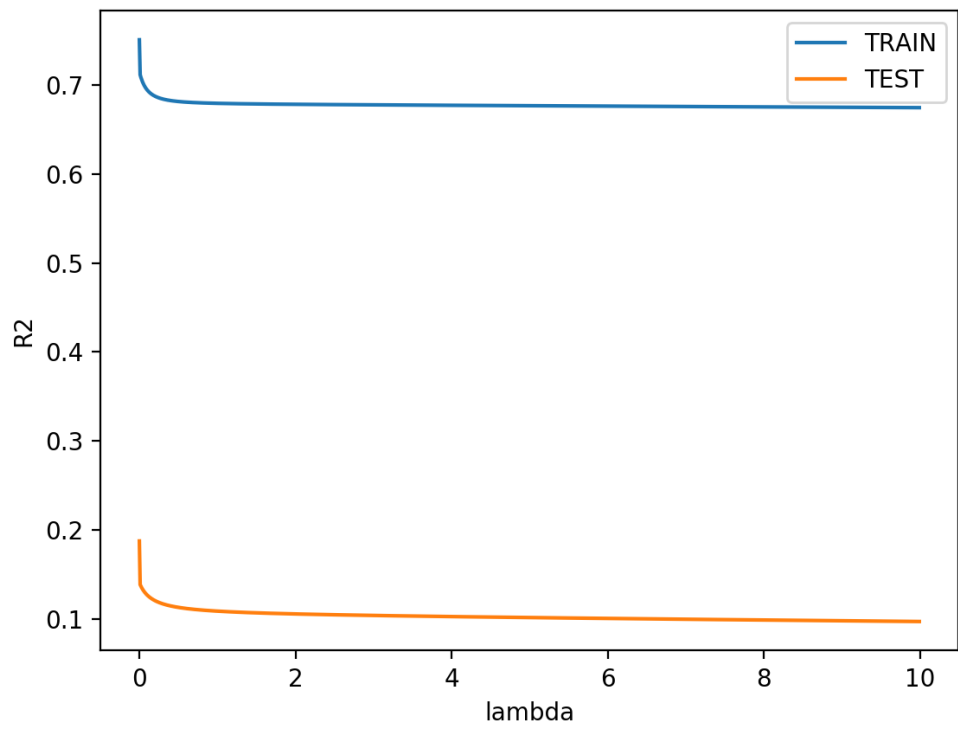
plt.plot(lamb_list, R2_2_list, label = 'TRAIN')
plt.plot(lamb_list, R2_test_2_list, label = 'TEST')
plt.xlabel('lambda')
plt.ylabel('R2')
plt.legend()
plt.show()

### Then the best lambda must be <1

### After so many tests

lamb_list = np.arange(0, 0.0001, 0.000001)
Theta2_list = []
R2_2_list = []
R2_test_2_list = []
for i in lamb_list:
    Theta2 = closed_form_2(X, Y, i)
    Theta2_list.append(Theta2)
    R2_2_list.append(R_square(np.dot(X, Theta2), Y))
    R2_test_2_list.append(R_square(np.dot(X_test, Theta2), Y_test))
'''

```



Problem 3:



(1)

```
print "Problem 3"

corr = df.iloc[:,2:10].corr()
f, ax = plt.subplots(figsize=(12, 10))
print(corr)
print(corr > 0.8)
```

```
Problem 3
```

	MEI	CO2	CH4	...	CFC-12	TSI	Aerosols
MEI	1.000000	-0.152911	-0.105555	...	-0.039836	-0.076826	0.352351
CO2	-0.152911	1.000000	0.872253	...	0.823210	0.017867	-0.369265
CH4	-0.105555	0.872253	1.000000	...	0.958237	0.146335	-0.290381
N2O	-0.162375	0.981135	0.894409	...	0.839295	0.039892	-0.353499
CFC-11	0.088171	0.401284	0.713504	...	0.831381	0.284629	-0.032302
CFC-12	-0.039836	0.823210	0.958237	...	1.000000	0.189270	-0.243785
TSI	-0.076826	0.017867	0.146335	...	0.189270	1.000000	0.083238
Aerosols	0.352351	-0.369265	-0.290381	...	-0.243785	0.083238	1.000000

```
[8 rows x 8 columns]
```

	MEI	CO2	CH4	N2O	CFC-11	CFC-12	TSI	Aerosols
MEI	True	False	False	False	False	False	False	False
CO2	False	True	True	True	False	True	False	False
CH4	False	True	True	True	False	True	False	False
N2O	False	True	True	True	False	True	False	False
CFC-11	False	False	False	False	True	True	False	False
CFC-12	False	True	True	True	True	True	False	False
TSI	False	False	False	False	False	False	True	False
Aerosols	False	False	False	False	False	False	False	True

```
[Finished in 5.0s]
```

We can find that CO2, CH4, N2O and CFC-12 are highly correlated, we can drop three of them, for example, CO2, CH4 and N2O.

(2)

```

127     print "Problem 3"
128
129     corr = df.iloc[:,2:10].corr()
130     f, ax = plt.subplots(figsize=(12, 10))
131     print(corr)
132     print(corr > 0.8)
133
134     new_df = df.copy()
135     new_df.drop(['C02', 'CH4', 'N20'], axis=1, inplace=True)
136     new_Training_set = new_df.loc[new_df['Year'] <= 2006]
137     new_Testing_set = new_df.loc[new_df['Year'] > 2006]
138     new_X = np.matrix(Training_set.iloc[:,2:11].values)
139     new_Y = np.matrix(Training_set.iloc[:,11].values).T
140     new_X_test = np.matrix(Testng_set.iloc[:,2:11].values)
141     new_Y_test = np.matrix(Testng_set.iloc[:,11].values).T
142
143     new_Theta = closed_form_2(new_X, new_Y, lamb)
144
145     print R_square(np.dot(new_X, new_Theta), new_Y)
146     print R_square(np.dot(new_X_test, new_Theta), new_Y_test)
147

```

Problem 3

	MEI	C02	CH4	...	CFC-12	TSI	Aerosols
MEI	1.000000	-0.152911	-0.105555	...	-0.039836	-0.076826	0.352351
C02	-0.152911	1.000000	0.872253	...	0.823210	0.017867	-0.369265
CH4	-0.105555	0.872253	1.000000	...	0.958237	0.146335	-0.290381
N20	-0.162375	0.981135	0.894409	...	0.839295	0.039892	-0.353499
CFC-11	0.088171	0.401284	0.713504	...	0.831381	0.284629	-0.032302
CFC-12	-0.039836	0.823210	0.958237	...	1.000000	0.189270	-0.243785
TSI	-0.076826	0.017867	0.146335	...	0.189270	1.000000	0.083238
Aerosols	0.352351	-0.369265	-0.290381	...	-0.243785	0.083238	1.000000

[8 rows x 8 columns]

	MEI	C02	CH4	N20	CFC-11	CFC-12	TSI	Aerosols
MEI	True	False	False	False	False	False	False	False
C02	False	True	True	True	False	True	False	False
CH4	False	True	True	True	False	True	False	False
N20	False	True	True	True	False	True	False	False
CFC-11	False	False	False	False	True	True	False	False
CFC-12	False	True	True	True	True	True	False	False
TSI	False	False	False	False	False	False	True	False
Aerosols	False	False	False	False	False	False	False	True

0.694468408539303

0.12660600610369788

[Finished in 2.7s]

Problem 4:

```
def Gradient_Descent(alpha, theta_0, x, y, tol):
    theta = theta_0
    cost = (1./len(x)) * (x.T @ (x @ theta_0 - y))
    count = 0
    while not ((np.all(cost) <= tol) | (count > 10000)):
        count += 1
        theta = theta - alpha * (1./len(x)) * (x.T @ (x @ theta - y))
        cost = (1./len(x)) * (x.T @ (x @ theta - y))
    print('iterate {} times'.format(count))
    return theta
```